

Rover vs. Splunk Enterprise Security

Keep security history within reach - not just in storage.

Compare five priorities: historical access, deployment, investigation context, detection execution and annual cost.

Scope: Splunk means Enterprise Security (ES), with cloud, self-managed and edition differences identified where relevant. ^{[2][10]} Descriptions reflect vendor materials, not independent tests. **Confirm** marks an evaluation item, not an unavailable feature.

01 Data retention & search

EVALUATION AREA	ROVER	SPLUNK ENTERPRISE SECURITY
Data architecture and ownership	Data and indexes in customer-owned object storage. ^[1]	Runs on Splunk-managed cloud or customer-managed infrastructure. SmartStore can use remote object storage with a local indexer cache. ^{[2][3]}
Multi-year retention	Multi-year searchable retention advertised; confirm the contracted duration. ^[1]	Retention is configurable. Distinguish directly searchable retention from archival retention and size the corresponding storage. ^[4]
Historical access and restoration	Search retained history without restoring archived logs. ^[1]	Access depends on the storage path. Managed Active Archive requires restoration; Federated Search for DDSS can query eligible Amazon S3 archives from Splunk Cloud deployments in AWS regions. ^{[5][6]}
Search infrastructure	Serverless query compute; no persistent search clusters. ^[1]	Search capacity and concurrency depend on platform resources and settings. Infrastructure responsibilities differ between managed cloud and self-managed deployments. ^{[2][13]}
Search and analytics	Full-text search, field filters and aggregations. ^[1]	SPL supports correlations, lookups, joins and analytical commands. Federated Search adds access to supported external data stores; deployment and source restrictions apply. ^{[8][6]}

BUYER QUESTION / Can we investigate a two-year-old incident without restoring logs?

Ask each vendor to demonstrate the exact storage and query path in its proposal. A retention commitment should specify the access steps, supported query functions and cost, not just the age of stored data.

Proof to request: The same historical dataset, query mix and concurrency. Record result completeness, search latency, access steps and incremental cost.

Deployment & ongoing operations

Measure the time to usable security coverage. First ingestion, production onboarding and complete migration are different milestones.

EVALUATION AREA	ROVER	SPLUNK ENTERPRISE SECURITY
Data ingestion methods	HTTP push/pull and object-storage ingestion. ^[1]	Forwarders, HTTP Event Collector, standard and custom inputs, and source-specific apps and add-ons. ^[8]
Initial onboarding	Advertises initial setup in hours, not a migration guarantee. ^[1]	Request a deployment-specific plan. Data-onboarding tools and implementation services are available; timing depends on data sources, detections and workflow scope. ^[8]
Schema and field preparation	Schema-on-read reduces upfront schema work. ^[1]	Supports search-time extraction, calculated fields, aliases and lookups. Search-time schema handling is not exclusive to Rover. ^[9]
Ongoing administration	No persistent search-cluster administration. Confirm remaining responsibilities. ^[1]	Splunk manages cloud-platform infrastructure; self-managed Enterprise adds customer infrastructure duties. Data inputs, searches and security content still need administration. ^{[2][8]}
Connector coverage and migration	Confirm: supported connectors, data formats, regions and migration assistance.	Splunkbase provides apps and add-ons for data sources and use cases. Validate compatibility with the selected platform, ES version and required integrations. ^[8]

BUYER QUESTION / Can we replace our existing SIEM in hours?

Treat initial setup as the start of an evaluation, not the completion of a migration. Agree on the integrations, detections, dashboards, access controls and response workflows that must work before cutover.

WHAT TO VALIDATE

Define production readiness. Demonstrate a working data source, a validated detection, an investigation and a response handoff.

Make responsibilities explicit. Assign owners for data quality, access, retention, upgrades, content tuning and operational support.

Plan a reversible migration. Inventory existing content; agree on parallel operation, historical-data access, rollback and exit criteria.

Evaluation checkpoint: compare staff effort and implementation scope for the same coverage, rather than comparing the first-log timestamp.

Investigation context & AI

Compare the evidence an analyst can inspect, the manual work required and the response actions the team can safely complete.

EVALUATION AREA	ROVER	SPLUNK ENTERPRISE SECURITY
Historical investigation context	Retained history and cross-source timelines support investigations. ^[1]	Analyst Queue and Threat Topology support investigation across findings, risk objects, assets and users. Historical access depends on the configured storage and search path. ^{[11][4][6]}
AI assistance	AI Context Engine uses retained security history. ^[1]	Advertises natural-language queries, investigation guidance, summaries and reports. Some AI capabilities have controlled availability; confirm cloud, version and feature eligibility. ^{[10][11]}
Behavioral analytics	Confirm: behavioral baselines, anomaly models and production availability.	UEBA is included in the published Premier bundle, not Essentials . Confirm the required capabilities in the proposed edition. ^[10]
Response automation and case management	Confirm: response integrations, playbooks, approvals and case-management capabilities.	Premier includes SOAR for orchestration and playbooks. ES also documents adaptive response actions and response plans; verify the edition and configuration for the required workflow. ^{[10][11]}

BUYER QUESTION / What should an AI-assisted investigation deliver beyond a summary?

Require a traceable explanation: the underlying events, entity links, search steps and timeline should be inspectable. Use the same incident to compare what analysts can verify and what they must investigate manually.

WHAT TO VALIDATE

Trace every important claim. Ask the analyst to open the original event behind each important conclusion and inspect its timestamp and source.

Measure analyst effort. Record manual pivots, query rewrites, missing context and the time to produce a reviewable incident record.

Validate safe response. Demonstrate approval gates, a ticket or case handoff, action logging and the required response integration.

Availability checkpoint: a feature listed on a product page is not a guarantee that it is enabled in every edition, region or deployment.

Detection coverage & execution

Start with the risks you need to detect. Then test whether the required rules remain timely and complete when telemetry and investigation demand rise.

EVALUATION AREA	ROVER	SPLUNK ENTERPRISE SECURITY
Detection execution model	Signal Mesh continuously updates detection state as telemetry arrives. ^[1]	Supports search-based detections over events and findings. Scheduled execution uses configurable platform scheduling controls. ^{[12][11][13]}
Detection versus investigation workloads	Detection and historical queries use separate execution paths. ^[1]	Scheduled workloads use concurrency controls. Supported versions also offer dynamic allocation between ad hoc and scheduled searches; confirm the proposed configuration. ^[13]
Behavior under heavy load	Validate: alert latency, backlog recovery and result completeness under representative load.	Real-time scheduling can skip runs during contention; continuous scheduling defers them. Configuration and capacity determine observed behavior. ^[13]
Detection content and coverage	Confirm: detection library, rule types, ATT&CK mappings and tuning tools.	ES Content Updates and the Use Case Library provide research-backed detections and ATT&CK mappings. Detection Studio supports testing, deployment and monitoring; verify availability. ^{[11][10]}

BUYER QUESTION / What happens when detection workloads and investigations peak together?

Run the detection set while replaying a telemetry burst and concurrent historical hunts. Compare detection delay, missing or duplicate findings, recovery time and the operational intervention required. Architecture alone does not establish a reliability advantage.

WHAT TO VALIDATE

- Keep coverage equivalent.** Use matching data sources, risk scenarios, detection logic, lookback periods and alert expectations.
- Capture contention and recovery.** Measure alert latency and completeness before, during and after the burst. Inspect scheduling or backlog telemetry.
- Test detection maintenance.** Change a rule, test it, deploy it and roll it back. Evaluate late-arriving data and rule-version behavior.

Terminology note: Splunk's "real-time scheduling" is a scheduler mode, not the same thing as a real-time search. Continuous scheduling prioritizes completeness over recency. ^[13]

Pricing & annual cost

Compare the annual cost of the same security coverage. A subscription price alone does not establish total cost or equivalent functionality.

EVALUATION AREA	ROVER	SPLUNK ENTERPRISE SECURITY
Primary pricing basis	Daily-ingestion-based pricing. ^[1]	Quote-based. Ingest, workload and activity-based options vary by offering and deployment; confirm the platform and ES bundle. ^[2]
Query charges	\$0.01 per query ; confirm billable-query definitions and limits. ^[1]	Query economics depend on the pricing model. Federated Search is priced by data volume scanned, rather than a published flat per-query rate. ^{[2][7]}
Retention and storage costs	Multi-year retention advertised as included; clarify underlying cloud charges. ^[1]	Storage arrangements vary. Request explicit pricing for searchable storage, archival storage and applicable external-storage or federated-search usage. ^{[2][7]}
Annual pricing at 1 TB/day	Request an annual quote for the agreed ingestion, retention and query workload.	Request a matched annual quote specifying ES edition, deployment, pricing model, retention and services. ^[2]
Total-cost evaluation	Include subscription, billable queries, applicable cloud costs, implementation, support and administration.	Include platform and ES licensing, applicable storage and federation, implementation, support and administration. Avoid double-counting capabilities included in a bundle.

BUYER QUESTION / Which platform is less expensive for our environment?

A defensible answer requires matched quotes and a representative workload. Specify ingestion, retention, query frequency, detection coverage, support and deployment responsibilities; then compare first-year and ongoing annual costs separately.

WHAT TO VALIDATE

Agree on billable units. Define raw versus processed ingestion, billable queries, scanned-data units, minimums, overages and exclusions.

Separate one-time and recurring cost. Show implementation and migration apart from subscription, cloud usage and ongoing administration.

Model growth and exit. Request the cost of higher volume, a longer retention period, additional investigations and data export at termination.

Commercial checkpoint: included retention does not, by itself, clarify who pays underlying object-storage, request, retrieval or data-transfer charges. Resolve these items in the quote.

EVALUATION WORKSHEET

Make the comparison measurable

Use one agreed scenario for both proposals. The example below is a planning baseline, not a vendor quote, benchmark or sizing guarantee.

COMMON INPUT	AGREED BASELINE
Ingestion	1 TB/day average, using decimal TB and the same pre/post-filtering definition. Document peak volume and expected growth.
Retention	Illustrative target: 3 years of searchable history. At a constant 1 TB/day, 365-day years imply 1,095 TB of raw input over three years, before compression, indexes or replicas.
Investigations	Agree on monthly query count, query mix, lookback periods and concurrent analysts or AI agents. For scan-based pricing, also measure bytes scanned.
Security coverage	Use the same data sources, detection use cases, response integrations and support level. Record any excluded or separately licensed capabilities.

Annual cost worksheet

Ask each vendor to complete the same categories in the same currency. Mark included items explicitly; do not add them twice.

COST CATEGORY	ROVER	SPLUNK ENTERPRISE SECURITY
Subscription / platform / SIEM	Vendor quote	Vendor quote
Query or federation usage	Workload-based estimate	Workload-based estimate
Applicable storage and cloud charges	Itemized estimate	Itemized estimate
Support and ongoing administration	Included or itemized	Included or itemized
Implementation, migration and training	One-time estimate	One-time estimate

First-year cost = recurring annual charges + one-time implementation and migration.
Ongoing annual cost = recurring charges at the agreed operating scale.

Query-fee illustration only: at 100,000 billable queries per month and Rover's published rate, annual query fees would be 100,000 x 12 x \$0.01 = \$12,000. This excludes the ingestion subscription and any other applicable charges. ^[1]

See Rover with your own security data

Bring one high-volume source, a historical investigation and your required detection workflows. Review the evidence, operating responsibilities and annual quote with Rover.

[Book a Rover demo](#) | [Rover website](#)

REFERENCE

Sources & comparison notes

Public vendor materials reviewed on 30 September 2026. Click a numbered citation in the comparison or a source title below to open the reference.

^[1] **Rover: product, architecture and pricing**

Rover public website; architecture, ingestion, Signal Mesh, search, AI context and published query rate.

^[2] **Splunk: pricing and purchasing options**

Platform deployment options, ES editions and commercial pricing models.

^[3] **Splunk Enterprise 10.4: About SmartStore**

Remote object storage, indexers and local caching.

^[4] **Splunk Cloud 10.5.2605: manage indexes**

Searchable retention, index settings and archive options.

^[5] **Splunk Cloud 10.5.2605: managed archive**

Dynamic Data Active Archive and restoration for search.

^[6] **Splunk Cloud 10.5.2605: federated-search options**

External data access; DDSS support and AWS-region restrictions.

^[7] **Splunk: Federated Search**

External data-lake access and pricing by data volume scanned.

^[8] **Splunk Enterprise: product features**

Ingestion, SPL, apps, add-ons, support and services.

^[9] **Splunk Enterprise 10.4: search-time operations**

Field extractions, aliases, calculated fields and lookups.

^[10] **Splunk Enterprise Security: product and editions**

AI availability and Essentials/Premier feature packaging.

^[11] **Splunk Enterprise Security: feature reference**

Analyst Queue, Threat Topology, response workflows and detection content.

^[12] **Splunk ES 8.5: event-based detections**

Search-based detections, trigger conditions and adaptive response.

^[13] **Splunk Enterprise 10.0: scheduled-search priority**

Concurrency, dynamic allocation and real-time versus continuous scheduling.

How to use this guide

Vendor descriptions are not independent performance, reliability or cost findings. The evaluation prompts and worksheet are proposed checks, not measured results. No overall winner, migration-time guarantee or matched annual-price advantage is established.

Feature availability and commercial terms may vary by version, edition, region and agreement. Splunk federation, archive restoration, SmartStore caching and continuous scheduling describe different mechanisms and should not be treated as interchangeable.

Before relying on a proposal, confirm supported integrations, contracted retention, AI and response features, deployment responsibilities, service levels and the treatment of underlying cloud charges. Product and company names identify their respective owners; this guide does not imply Splunk endorsement.